

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score
from sklearn.cluster import KMeans
from statsmodels.tsa.arima.model import ARIMA
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Load the dataset
```

```
# drive.mount('/content/drive')
```

```
# file_path = '/content/drive/MyDrive/healthcare_dataset.csv'
```

```
file_path = "healthcare_ds.csv"
df = pd.read_csv(file_path)
```

```
# Display the shape of the data and first few rows
```

```
print(f"Dataset shape: {df.shape}")
df.head()
```

```
Dataset shape: (55500, 15)
```

	Name	Age	Gender	Blood Type	Medical Condition	Date of Admission \
0	Bobby JacksOn	30	Male	B-	Cancer	2024-01-31
1	LesLie TErRy	62	Male	A+	Obesity	2019-08-20
2	DaNnY sMith	76	Female	A-	Obesity	2022-09-22
3	andrEw waTtS	28	Female	O+	Diabetes	2020-11-18
4	adRIENNE bEll	43	Female	AB+	Cancer	2022-09-19

	Doctor	Hospital	Insurance Provider \
0	Matthew Smith	Sons and Miller	Blue Cross
1	Samantha Davies	Kim Inc	Medicare
2	Tiffany Mitchell	Cook PLC	Aetna
3	Kevin Wells	Hernandez Rogers and Vang,	Medicare
4	Kathleen Hanna	White-White	Aetna

	Billing Amount	Room Number	Admission Type	Discharge Date
0	18856.281306	328	Urgent	2024-02-02

Paracetamol

1	33643.327287	265	Emergency	2019-08-26
Ibuprofen				
2	27955.096079	205	Emergency	2022-10-07
Aspirin				
3	37909.782410	450	Elective	2020-12-18
Ibuprofen				
4	14238.317814	458	Urgent	2022-10-09
Penicillin				
Test Results				
0	Normal			
1	Inconclusive			
2	Normal			
3	Abnormal			
4	Abnormal			

Drop unnecessary columns (Name, Doctor, Hospital), encode categorical variables, and handle missing data

```
# Drop columns that are not needed for ML models
df_cleaned = df.drop(['Name', 'Doctor', 'Hospital'], axis=1)

# Convert date columns to datetime and create "Length of Stay"
df_cleaned['Date of Admission'] = pd.to_datetime(df_cleaned['Date of Admission'])
df_cleaned['Discharge Date'] = pd.to_datetime(df_cleaned['Discharge Date'])
df_cleaned['Length of Stay'] = (df_cleaned['Discharge Date'] - df_cleaned['Date of Admission']).dt.days

# Handle missing values (e.g., replace NaN with mean or mode)
df_cleaned['Billing Amount'].fillna(df_cleaned['Billing Amount'].mean(), inplace=True)
df_cleaned['Medication'].fillna('Unknown', inplace=True)
df_cleaned['Test Results'].fillna('Unknown', inplace=True)

# Encode categorical variables
label_encoder = LabelEncoder()

# Encode 'Gender', 'Blood Type', 'Medical Condition', 'Admission Type', 'Test Results'
df_cleaned['Gender'] = label_encoder.fit_transform(df_cleaned['Gender'])
df_cleaned['Blood Type'] = label_encoder.fit_transform(df_cleaned['Blood Type'])
df_cleaned['Medical Condition'] = label_encoder.fit_transform(df_cleaned['Medical Condition'])
df_cleaned['Admission Type'] = label_encoder.fit_transform(df_cleaned['Admission Type'])
```

```
df_cleaned['Test Results'] =  
label_encoder.fit_transform(df_cleaned['Test Results'])
```

```
# Display first few rows after preprocessing  
df_cleaned.head()
```

C:\Users\roshan\AppData\Local\Temp\ipykernel_1600\2971295002.py:10:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

```
df_cleaned['Billing Amount'].fillna(df_cleaned['Billing  
Amount'].mean(), inplace=True)
```

C:\Users\roshan\AppData\Local\Temp\ipykernel_1600\2971295002.py:11:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

```
df_cleaned['Medication'].fillna('Unknown', inplace=True)
```

C:\Users\roshan\AppData\Local\Temp\ipykernel_1600\2971295002.py:12:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

```
df_cleaned['Test Results'].fillna('Unknown', inplace=True)
```

	Age	Gender	Blood Type	Medical Condition	Date of Admission	\
0	30	1	5	2	2024-01-31	
1	62	1	0	5	2019-08-20	
2	76	0	1	5	2022-09-22	
3	28	0	6	3	2020-11-18	
4	43	0	2	2	2022-09-19	

	Insurance Provider	Billing Amount	Room Number	Admission Type	\
0	Blue Cross	18856.281306	328	2	
1	Medicare	33643.327287	265	1	
2	Aetna	27955.096079	205	1	
3	Medicare	37909.782410	450	0	
4	Aetna	14238.317814	458	2	

	Discharge Date	Medication	Test Results	Length of Stay
0	2024-02-02	Paracetamol	2	2
1	2019-08-26	Ibuprofen	1	6
2	2022-10-07	Aspirin	2	15
3	2020-12-18	Ibuprofen	0	30
4	2022-10-09	Penicillin	0	20

Predict 'Medical Condition' using Random Forest Classifier

Features (X) and Target (y)

```
X = df_cleaned.drop(['Medical Condition'], axis=1)
```

```
y = df_cleaned['Medical Condition']
```

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Initialize RandomForestClassifier

```
rf_classifier = RandomForestClassifier(n_estimators=100,
random_state=42)
```

Convert 'Date of Admission' and 'Discharge Date' to numerical features

```
X_train['Date of Admission'] = X_train['Date of
Admission'].dt.dayofyear
```

```
X_train['Discharge Date'] = X_train['Discharge Date'].dt.dayofyear
```

```
X_test['Date of Admission'] = X_test['Date of Admission'].dt.dayofyear
```

```
X_test['Discharge Date'] = X_test['Discharge Date'].dt.dayofyear
```

Convert 'Insurance Provider' and 'Medication' to numerical using Label Encoding

Create a LabelEncoder instance

```
insurance_encoder = LabelEncoder()
```

```
medication_encoder = LabelEncoder() # Create a new encoder for
'Medication'
```

```
# Fit and transform on training data
X_train['Insurance Provider'] =
insurance_encoder.fit_transform(X_train['Insurance Provider'])
X_train['Medication'] =
medication_encoder.fit_transform(X_train['Medication']) # Encode
'Medication'
```

```
# Transform test data using the same encoders
X_test['Insurance Provider'] =
insurance_encoder.transform(X_test['Insurance Provider'])
X_test['Medication'] =
medication_encoder.transform(X_test['Medication']) # Encode
'Medication'
```

```
# Train the model
rf_classifier.fit(X_train, y_train)

# Predict on test set
y_pred = rf_classifier.predict(X_test)
```

```
# Evaluate the model
print(f"Accuracy: {accuracy_score(y_test, y_pred)}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.291981981981982

Classification Report:

	precision	recall	f1-score	support
0	0.30	0.30	0.30	1915
1	0.29	0.30	0.29	1847
2	0.28	0.29	0.29	1871
3	0.29	0.29	0.29	1822
4	0.29	0.28	0.28	1788
5	0.32	0.29	0.30	1857
accuracy			0.29	11100
macro avg	0.29	0.29	0.29	11100
weighted avg	0.29	0.29	0.29	11100

```
# Use KMeans clustering to segment patients into clusters
kmeans = KMeans(n_clusters=3, random_state=42)
```

```
# Instead of converting 'Date of Admission' and 'Discharge Date'
again,
# create a copy of df_cleaned with the necessary preprocessing for
KMeans
# Exclude 'Cluster' from drop as it doesn't exist yet. It will be
```

```
created later
X_kmeans = df_cleaned.drop(['Medical Condition'], axis=1)


# Convert 'Date of Admission' and 'Discharge Date' to numerical
features
X_kmeans['Date of Admission'] = X_kmeans['Date of
Admission'].dt.dayofyear
X_kmeans['Discharge Date'] = X_kmeans['Discharge Date'].dt.dayofyear

# Convert 'Insurance Provider' and 'Medication' to numerical using
Label Encoding
# Create a LabelEncoder instance for 'Insurance Provider'
insurance_encoder = LabelEncoder()
# Create a LabelEncoder instance for 'Medication'
medication_encoder = LabelEncoder()

# Fit and transform on X_kmeans
X_kmeans['Insurance Provider'] =
insurance_encoder.fit_transform(X_kmeans['Insurance Provider'])
X_kmeans['Medication'] =
medication_encoder.fit_transform(X_kmeans['Medication'])

# Fit the KMeans model
df_cleaned['Cluster'] = kmeans.fit_predict(X_kmeans)

# Visualize clusters using Age and Billing Amount
plt.figure(figsize=(10,6))
sns.scatterplot(x='Age', y='Billing Amount', hue='Cluster',
data=df_cleaned, palette='viridis')
plt.title('Patient Clusters based on Age and Billing Amount')
plt.show()
```

Step 5: Resource Forecasting (Time-Series Analysis)

We'll use ARIMA to forecast hospital resource demand (e.g., room occupancy).

```
# Forecasting: Use 'Length of Stay' and 'Billing Amount' for time-
series prediction

# Resample the data by month for hospital resource demand prediction
df_cleaned['Date of Admission'] = pd.to_datetime(df_cleaned['Date of
Admission'])
df_monthly = df_cleaned.resample('ME', on='Date of
Admission').agg({'Billing Amount': 'sum'})

# Plot the total billing amount over time (as an indicator of hospital
resource use)
plt.figure(figsize=(10,6))
df_monthly['Billing Amount'].plot()
plt.title('Monthly Billing Amount (Hospital Resource Demand)')
plt.xlabel('Date')
plt.ylabel('Billing Amount')
plt.show()

# Fit ARIMA model (order may need tuning based on data
characteristics)
arima_model = ARIMA(df_monthly['Billing Amount'], order=(5, 1, 0)) #
Example order
```

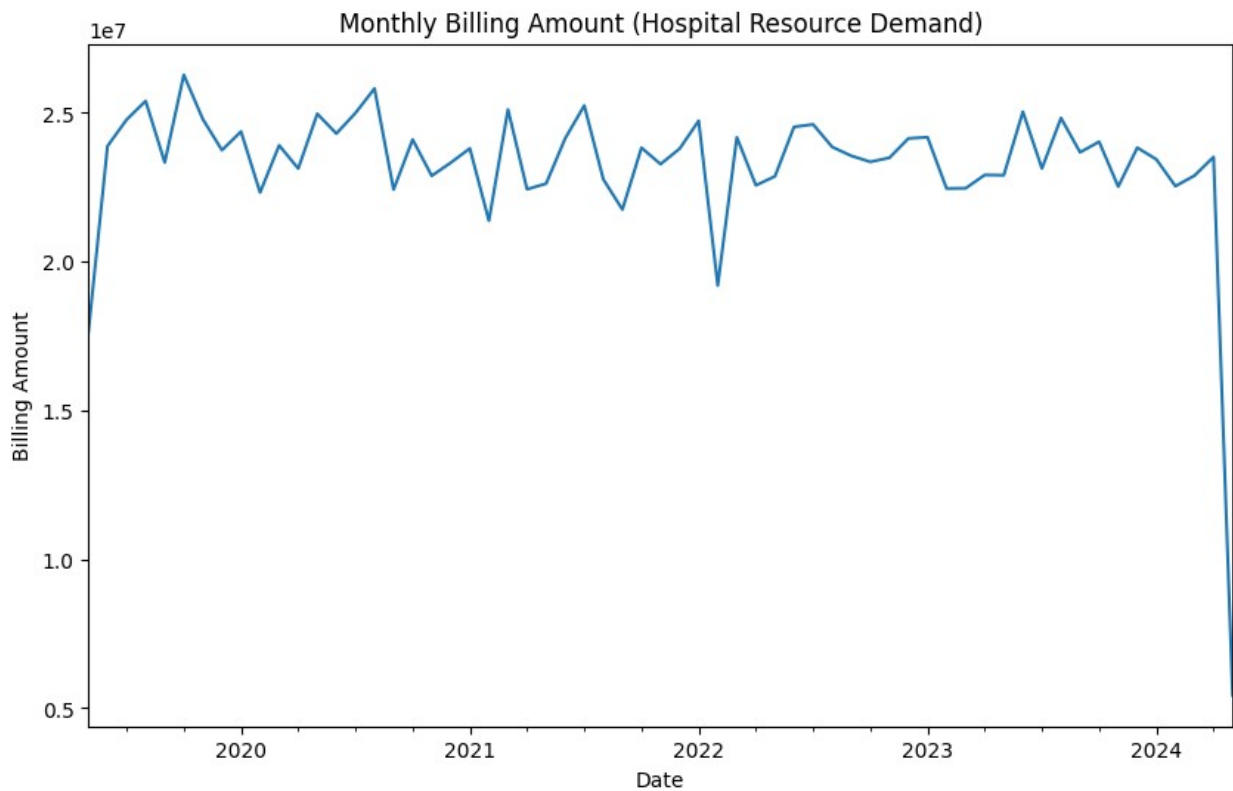
```


arima_fit = arima_model.fit()

# Forecast next 12 months
forecast = arima_fit.forecast(steps=12)
forecast_index = pd.date_range(df_monthly.index[-1], periods=12,
                                freq='ME')

# Plot forecast
plt.figure(figsize=(10,6))
plt.plot(df_monthly.index, df_monthly['Billing Amount'],
         label='Historical')
plt.plot(forecast_index, forecast, label='Forecast', color='red')
plt.title('Hospital Resource Forecast (Billing Amount)')
plt.xlabel('Date')
plt.ylabel('Billing Amount')
plt.legend()
plt.show()

```





Step 7: Model Explainability with SHAP For interpretability, SHAP (Shapley Additive Explanations) helps us understand how each feature contributes to the model's predictions. We'll focus on the RandomForestClassifier.


For interpretability, SHAP (Shapley Additive Explanations) helps us understand how each feature contributes to the model's predictions. We'll focus on the RandomForestClassifier.

```
import shap
from sklearn.datasets import load_iris
from sklearn.ensemble import RandomForestClassifier

# Load sample data
X, y = load_iris(return_X_y=True)
rf = RandomForestClassifier().fit(X, y)

# SHAP explanation
explainer = shap.TreeExplainer(rf)
shap_values = explainer.shap_values(X[:10]) # Small sample
shap.summary_plot(shap_values, X[:10])

<Figure size 640x480 with 0 Axes>
```



Step 8: Hyperparameter Tuning for RandomForestClassifier

You can further improve the Random Forest Classifier by tuning the hyperparameters using GridSearchCV.

```
from sklearn.model_selection import GridSearchCV

# Define the parameter grid for RandomForest
param_grid = {
    'n_estimators': [50, 100], # Reduced the n_estimators options
    'max_depth': [10, 20],     # Reduced the max_depth options
    'min_samples_split': [2, 5] # Reduced the min_samples_split
options
}

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=rf_classifier,
param_grid=param_grid, cv=3, n_jobs=-1, verbose=2)

# Fit the grid search to the data
grid_search.fit(X_train, y_train)

# Best parameters from GridSearchCV
print(f"Best parameters: {grid_search.best_params_}")

# Evaluate the best model on test data
best_rf_classifier = grid_search.best_estimator_
y_pred_best = best_rf_classifier.predict(X_test)
```

```
# Evaluate performance of the tuned model
print(f"Accuracy of tuned model: {accuracy_score(y_test,
y_pred_best)}")
print("Classification Report for tuned model:")
print(classification_report(y_test, y_pred_best))
```

Fitting 3 folds for each of 8 candidates, totalling 24 fits
Best parameters: {'max_depth': 20, 'min_samples_split': 5, 'n_estimators': 100}
Accuracy of tuned model: 0.2942342342342342
Classification Report for tuned model:

	precision	recall	f1-score	support
0	0.30	0.29	0.29	1915
1	0.31	0.29	0.30	1847
2	0.29	0.27	0.28	1871
3	0.26	0.29	0.27	1822
4	0.30	0.33	0.31	1788
5	0.31	0.29	0.30	1857
accuracy			0.29	11100
macro avg	0.29	0.29	0.29	11100
weighted avg	0.30	0.29	0.29	11100

Step 9: Final Model Evaluation

After tuning the model, you should assess the performance again using different metrics and analyze the results. For instance, you can use Confusion Matrix to see how well the model performs.

```
from sklearn.metrics import confusion_matrix
import seaborn as sns

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred_best)

# Plot confusion matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Class 0', 'Class 1', 'Class 2'], yticklabels=['Class 0', 'Class 1', 'Class 2'])
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```

