

```

# Install if needed
# !pip install shap imbalanced-learn

import pandas as pd
import numpy as np
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from statsmodels.tsa.arima.model import ARIMA
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score, davies_bouldin_score
from sklearn.metrics import classification_report, accuracy_score,
roc_auc_score, confusion_matrix, mean_absolute_error,
mean_squared_error
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt
import seaborn as sns

# Load dataset
df = pd.read_csv('healthcare_ds.csv')
df.head(5)

# Keep relevant columns only
df = df[['Age', 'Gender', 'Blood Type', 'Medical Condition',
        'Admission Type',
        'Test Results', 'Medication', 'Billing Amount', 'Date of
        Admission', 'Discharge Date']]

# Convert date columns
df['Date of Admission'] = pd.to_datetime(df['Date of Admission'],
errors='coerce')
df['Discharge Date'] = pd.to_datetime(df['Discharge Date'],
errors='coerce')
df['Length of Stay'] = (df['Discharge Date'] - df['Date of
Admission']).dt.days
df.dropna(inplace=True)

# Emergency flag
df['Is_Emergency'] = df['Admission Type'].apply(lambda x: 1 if x ==
'Emergency' else 0)

# Dividing diseases into Critical and Non-Critical
def categorize_condition(condition):
    condition = str(condition).lower()
    critical_keywords = ['cancer', 'tumor', 'stroke', 'cardiac',
'arrest', 'heart', 'failure', 'trauma', 'critical', 'coma']
    for keyword in critical_keywords:
        if keyword in condition:

```

```

        return 1
    return 0

df['Medical Condition'] = df['Medical
Condition'].apply(categorize_condition)

# Label Encoding
le = LabelEncoder()
for col in ['Gender', 'Blood Type', 'Admission Type', 'Test Results',
'Medication']:
    df[col] = le.fit_transform(df[col].astype(str))

df['Billing_per_day'] = df['Billing Amount'] / (df['Length of Stay'] +
1)

# Scale numerical features
scaler = StandardScaler()
df[['Age', 'Billing Amount', 'Length of Stay', 'Billing_per_day']] =
scaler.fit_transform(
    df[['Age', 'Billing Amount', 'Length of Stay', 'Billing_per_day']]
)

print("\n--- Disease Prediction (Random Forest) ---")
# Prepare features and target
X = df.drop(['Medical Condition', 'Date of Admission', 'Discharge
Date'], axis=1)
y = df['Medical Condition']

# Stratified split to preserve class balance
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)

print("Class distribution before SMOTE:")
print(y_train.value_counts())

#Apply SMOTE to balance the dataset
sm = SMOTE(random_state=42)
X_res, y_res = sm.fit_resample(X_train, y_train)

print("Class distribution after SMOTE:")
print(pd.Series(y_res).value_counts())

--- Disease Prediction (Random Forest) ---
Class distribution before SMOTE:
Medical Condition
0    37018
1     7382
Name: count, dtype: int64

```

Class distribution after SMOTE:

Medical Condition

0 37018

1 37018

Name: count, dtype: int64

#Train Random Forest Classifier

print("\n--- Training Random Forest Classifier ---")

rf = RandomForestClassifier(n_estimators=100, random_state=42)

rf.fit(X_res, y_res)

y_pred = rf.predict(X_test)

print("\n--- Binary Classification: Critical vs Non-Critical ---")

print("Accuracy Using RandomForestClassifier:", accuracy_score(y_test, y_pred))

print(classification_report(y_test, y_pred))

AUC-ROC Score

y_proba = rf.predict_proba(X_test)[: , 1]

roc_auc = roc_auc_score(y_test, y_proba)

print("AUC-ROC Score:", roc_auc)

#Train KNeighbors Classifier

print("\n--- Training KNeighborsClassifier ---")

knn = KNeighborsClassifier(n_neighbors=5)

knn.fit(X_res, y_res)

knn_pred = knn.predict(X_test)

knn_acc = accuracy_score(y_test, knn_pred)

print("Accuracy Using KNeighborsClassifier:", accuracy_score(y_test, knn_pred))

print(classification_report(y_test, knn_pred))

AUC-ROC Score

knn_proba = knn.predict_proba(X_test)[: , 1]

roc_auc = roc_auc_score(y_test, knn_proba)

print("AUC-ROC Score:", roc_auc)

cm_rfc = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(8,6))

sns.heatmap(cm_rfc, annot=True, fmt="d", cmap="coolwarm")

plt.title('Confusion Matrix (Random Forest Classifier)')

plt.ylabel('Actual Class')

plt.xlabel('Predicted Class')

plt.show()

cm_knc = confusion_matrix(y_test, knn_pred)

plt.figure(figsize=(8,6))

sns.heatmap(cm_knc, annot=True, fmt="d", cmap="coolwarm")

```
plt.title('Confusion Matrix (KNeighbors Classifier)')
plt.ylabel('Actual Class')
plt.xlabel('Predicted Class')
plt.show()
```

--- Training Random Forest Classifier ---

--- Binary Classification: Critical vs Non-Critical ---

Accuracy Using RandomForestClassifier: 0.7874774774774774

	precision	recall	f1-score	support
0	0.86	0.89	0.87	9255
1	0.33	0.26	0.29	1845
accuracy			0.79	11100
macro avg	0.59	0.58	0.58	11100
weighted avg	0.77	0.79	0.78	11100

AUC-ROC Score: 0.6230656248215642

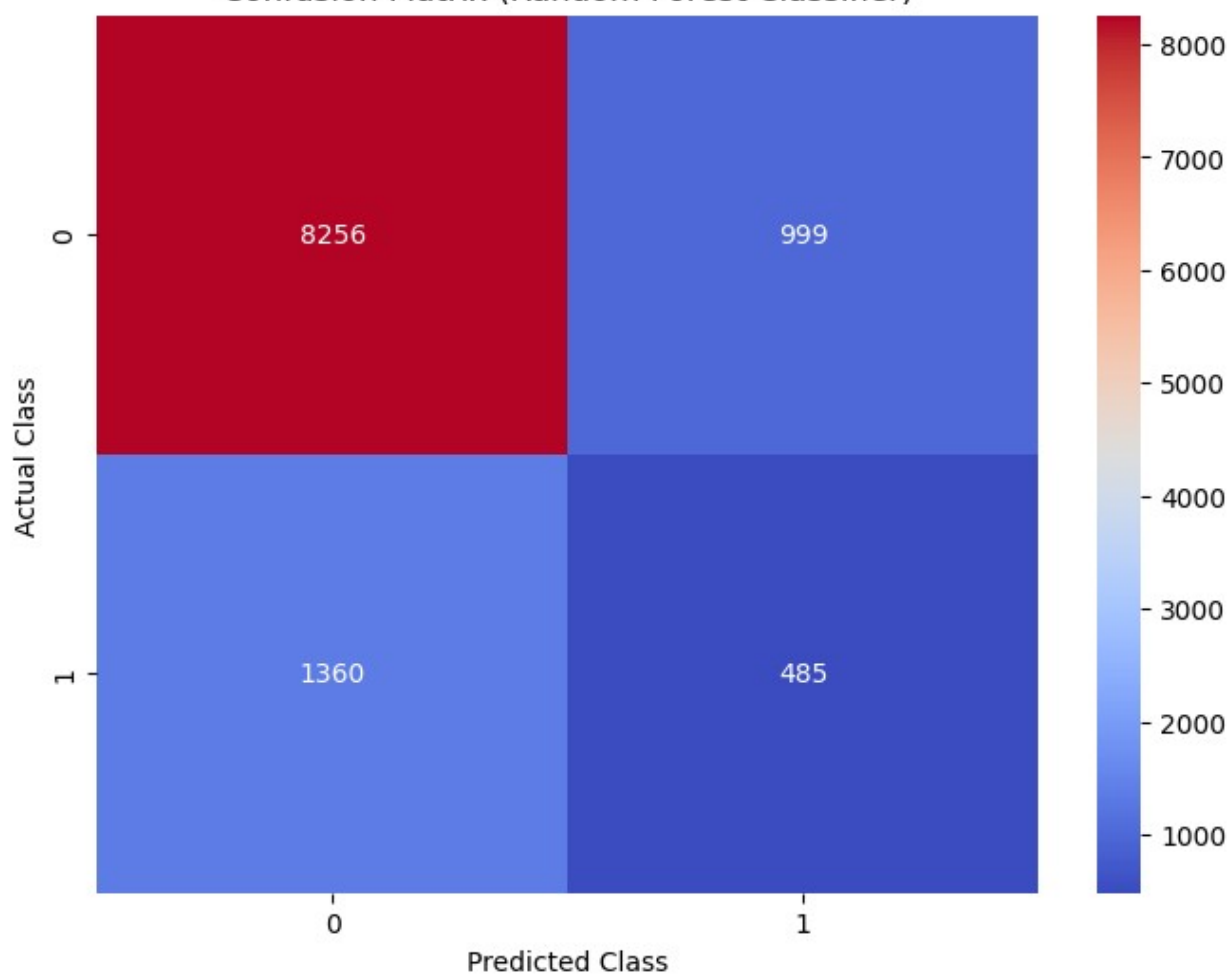
--- Training KNeighborsClassifier ---

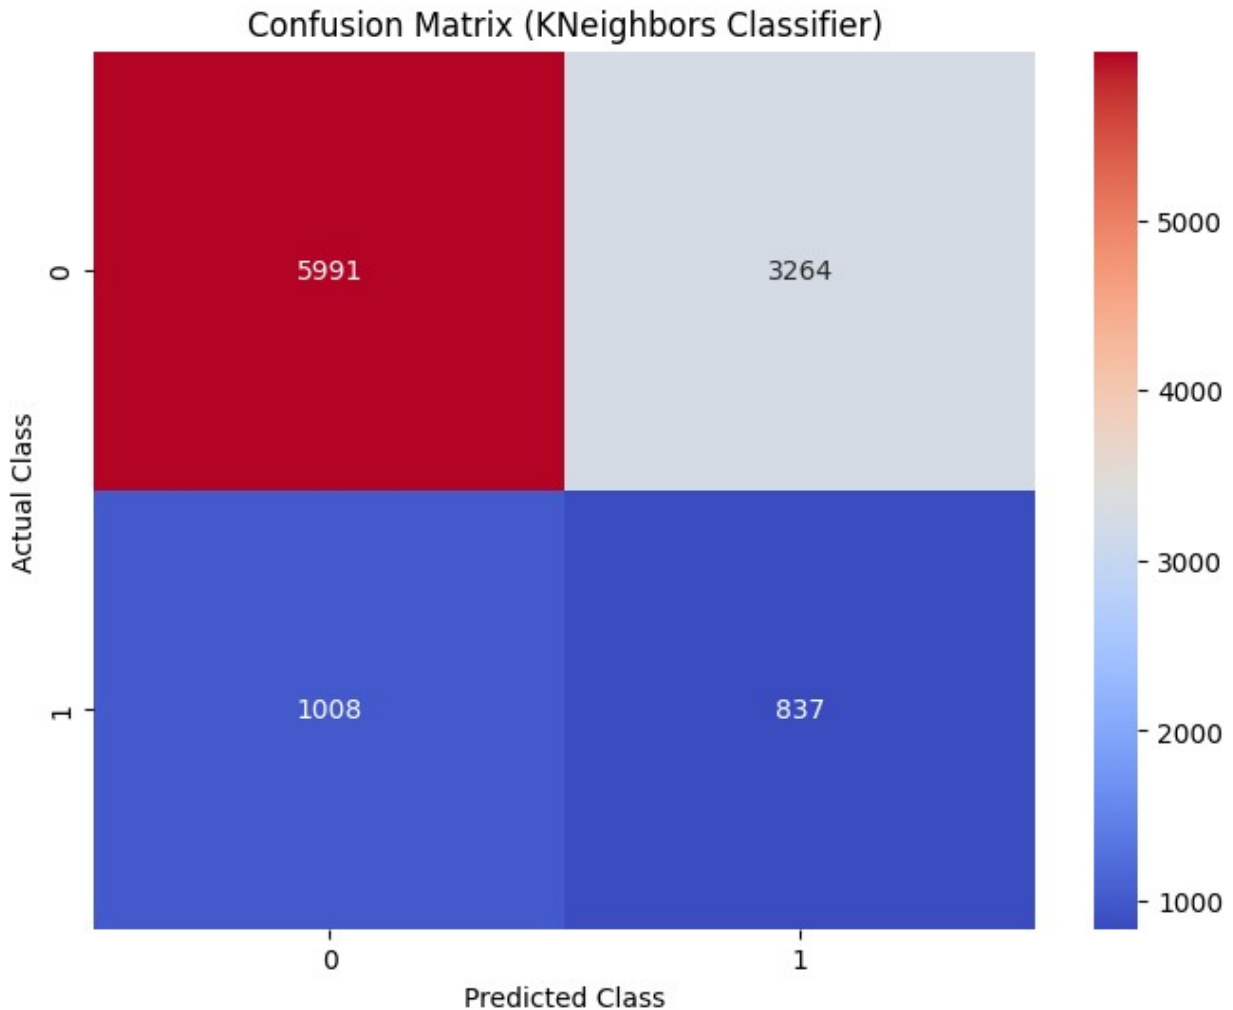
Accuracy Using KNeighborsClassifier: 0.6151351351351352

	precision	recall	f1-score	support
0	0.86	0.65	0.74	9255
1	0.20	0.45	0.28	1845
accuracy			0.62	11100
macro avg	0.53	0.55	0.51	11100
weighted avg	0.75	0.62	0.66	11100

AUC-ROC Score: 0.5639817633184436

Confusion Matrix (Random Forest Classifier)





```
from sklearn.cluster import AgglomerativeClustering

print("\n--- Patient Risk Clustering (K-Means) ---")
# Select and scale features
# df = df.sample(n=4000, random_state=42)
print(len(df))
cluster_data = df[['Age', 'Billing Amount', 'Length of Stay']]
scaler = StandardScaler()
scaled_data = scaler.fit_transform(cluster_data)

# Cluster patients using KMeans
kmeans = KMeans(n_clusters=3, random_state=42)
df['Risk Cluster'] = kmeans.fit_predict(scaled_data)

kmeans_sil = silhouette_score(scaled_data, df['Risk Cluster'])
db_score = davies_bouldin_score(scaled_data, df['Risk Cluster'])
print(f"Silhouette Score: {kmeans_sil:.3f}")
```

```

print(f"Davies-Bouldin Index: {db_score:.3f}")

agg = AgglomerativeClustering(n_clusters=3)
agg_labels = agg.fit_predict(scaled_data)
df['Agglo Cluster'] = agg_labels
agg_sil = silhouette_score(scaled_data, agg_labels)
agg_db = davies_bouldin_score(scaled_data, agg_labels)

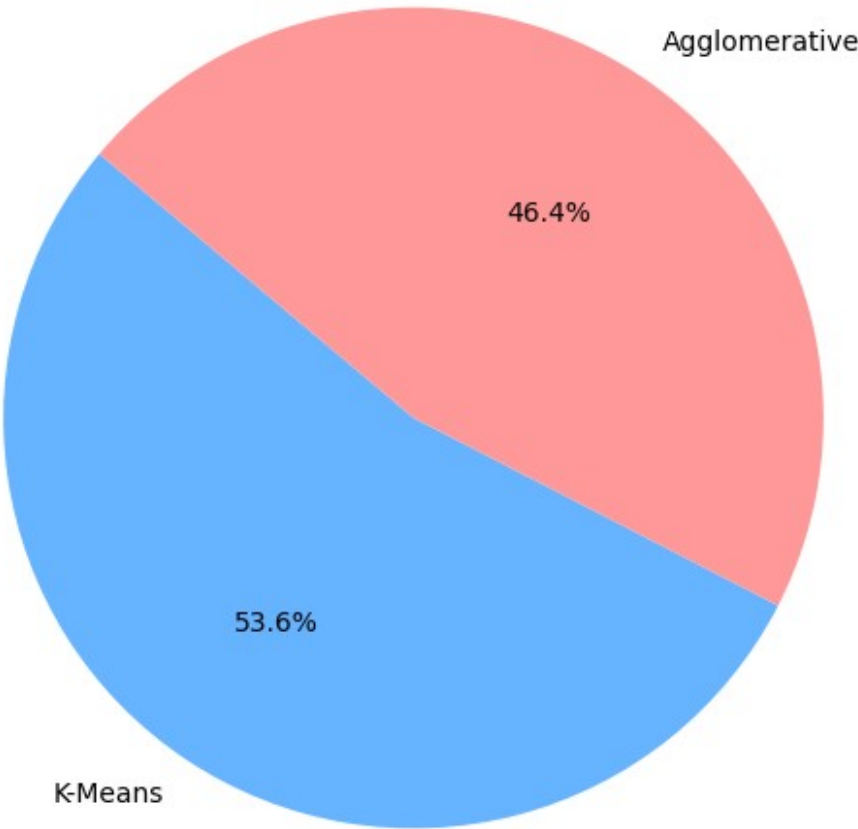
labels = ['K-Means', 'Agglomerative']
silhouette_scores = [kmeans_sil, agg_sil]
colors = ['#66b3ff', '#ff9999']
plt.figure(figsize=(6,6))
plt.pie(silhouette_scores, labels=labels, autopct='%1.1f%%',
        colors=colors, startangle=140)
plt.title('Clustering Model Comparison (Silhouette Score)')
plt.axis('equal')
plt.show()

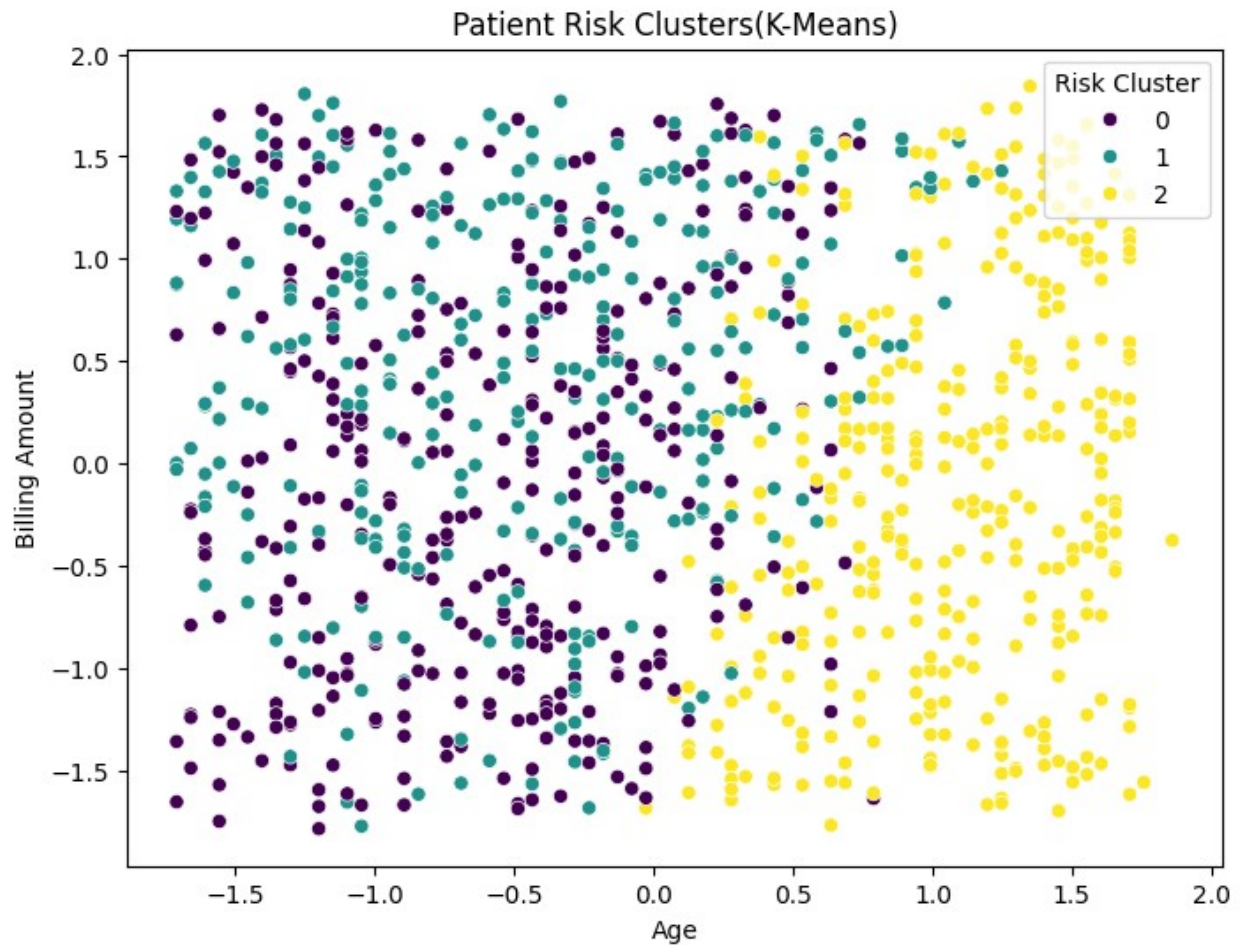
# Visualize KMeans clusters
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='Age', y='Billing Amount', hue='Risk
Cluster', palette='viridis')
plt.title('Patient Risk Clusters(K-Means)')
plt.show()

--- Patient Risk Clustering (K-Means) ---
1000
Silhouette Score: 0.248
Davies-Bouldin Index: 1.306

```

Clustering Model Comparison (Silhouette Score)






```
print("\n--- Hospital Resource Forecasting (ARIMA) ---")
admissions = df.groupby('Date of
Admission').size().asfreq('D').fillna(0)
model = ARIMA(admissions, order=(3, 1, 2))
model_fit = model.fit()
forecast = model_fit.forecast(steps=30)

plt.figure(figsize=(10, 5))
admissions.plot(label='Observed')
forecast.plot(label='Forecast', color='red')
plt.legend()
plt.title('Hospital Admissions Forecast')
plt.ylabel('Admissions Count')
plt.show()

--- Hospital Resource Forecasting (ARIMA) ---
```

```
c:\Users\roshan\Desktop\Dissertation\.venv\Lib\site-packages\
statsmodels\tsa\statespace\sarimax.py:966: UserWarning: Non-stationary
starting autoregressive parameters found. Using zeros as starting
parameters.
```

```
warn('Non-stationary starting autoregressive parameters')
c:\Users\roshan\Desktop\Dissertation\.venv\Lib\site-packages\
statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible
starting MA parameters found. Using zeros as starting parameters.
warn('Non-invertible starting MA parameters found.')
```



```
predicted = model_fit.predict(start=len(admissions)-30,
end=len(admissions)-1, dynamic=False)
true = admissions[-30:]
print("MAE:", mean_absolute_error(true, predicted))
print("RMSE:", np.sqrt(mean_squared_error(true, predicted)))
```

```
MAE: 4.467349574327432
RMSE: 5.662922822234962
```

```
# # ----- Step 8: SHAP Explainability -----
# explainer = shap.TreeExplainer(rf)
# shap_values = explainer.shap_values(X_test)
# shap.summary_plot(shap_values, X_test, plot_type="bar")
```